

Data Card

Input vs. Output	• What are users inputting (for example, audio, text, images)? • What are the outputs or labels your model will predict or generate?	Your response: Users will input natural-language text containing general health and medical questions. The dataset itself will contain medical documents, document chunks, question answer examples, and associated metadata. The system will generate natural-language responses to health-related questions, ideally grounded in relevant medical sources. Dataset records may also contain metadata such as medical topic, source, publication date, risk level, and whether a disclaimer is required.
Classification problem	• What is the classification task, if any? • What categories or classes are you working with, and why?	Your response: Classification is not the primary task of MedAssist. However, classification may be used to support retrieval and safety mechanisms. Potential categories include medical topic (e.g., general health, symptoms, diseases, prevention, medication information, nutrition, maternal health, mental health) and risk level (low, moderate, or high risk). These categories will help determine how information should be retrieved and

		presented and when additional warnings, escalation, or refusal may be appropriate. They will not be used to diagnose users.
Labeling	• Where do your labels come from—domain experts, community members, crowd workers, AI-generated sources? • How will you check for consistency and accuracy?	Your response: Labels will primarily be derived from trusted medical sources and predefined annotation guidelines. Where feasible, team members with relevant medical or domain knowledge will review high-risk content and a sample of labeled records. AI tools may assist with preliminary classification, but AI-generated labels will not automatically be treated as ground truth. We will manually review samples for inconsistencies, compare records against their original sources, and refine our labeling guidelines when disagreements or errors are identified.
Data creation	• How will you collect your data? • Will it come from direct sources (like interviews or recordings), or will you synthesize data in low-resource settings (e.g., for under-resourced languages)?	Your response: Our data will primarily be collected and structured from authoritative, publicly available medical sources, including clinical guidelines, public health organizations, peer-reviewed literature, and reputable medical reference resources. We will create structured document chunks for retrieval and evidence-based question-answer

		examples for evaluation and, if justified by later experiments, potential fine-tuning. We do not currently plan to collect direct patient interviews or recordings. We will also avoid using identifiable patient records. If synthetic question–answer examples are created to address gaps, they will be clearly identified and reviewed against authoritative medical sources.
Web collection ethics	• Do you plan to source data from the web? • If yes, which sites or platforms? • What tools will you use - AI tools, human workers, or both? • Have you addressed issues of consent, copyright, and privacy?	Your response: Yes. We plan to source information from authoritative medical and public health websites, such as WHO, CDC, national health agencies, clinical guideline repositories, and publicly accessible peer-reviewed medical literature. We may use both automated and manual methods for collection, with AI tools assisting with cleaning, structuring, or preliminary classification. Human review will remain part of quality control. We will consider the terms of use, licensing, copyright, attribution requirements, and privacy implications of each source. Public availability will not be treated as permission to collect personal health information. We will avoid identifiable patient

		information and will not intentionally collect private or sensitive health data.
Bias and representation	• Who might be left out of your dataset? • How will you ensure that the dataset reflects the diversity of the community and includes underrepresented voices?	Your response: The dataset may underrepresent people from low-resource healthcare settings, populations poorly represented in major medical research, speakers of underrepresented languages, and communities whose healthcare experiences differ from those represented in major international medical literature. Medical evidence may also be disproportionately based on populations and healthcare systems from high-income countries, while English-language sources may dominate the dataset. We will seek geographically and contextually diverse medical sources where they meet appropriate quality and evidence standards. We will document known representation gaps rather than claiming that the dataset represents every population equally. We will not introduce unsupported medical claims simply to increase representation.

Values and community alignment

- How does your dataset design reflect the values you identified in [01 Build Your Own Small Language Model](#)
- How will your community benefit from this dataset, and how will you ensure accountability to them?

Your response:
Our dataset design primarily reflects the values of scientific rationalism, universalism, and liberty identified in Course 1.

Scientific rationalism guides our reliance on evidence, clinical guidelines, peer-reviewed research, and verifiable medical sources. Universalism guides our goal of making reliable health information more accessible across different socioeconomic and geographic contexts. Liberty guides our commitment to empowering people with understandable information so they can make more informed health decisions while understanding the limitations of AI.

We will also prioritize privacy, safety, transparency, fairness, and accountability. We will document our data sources, collection methods, filtering decisions, labeling process, limitations, and known representation gaps so that our decisions can be reviewed and improved.